

Implement:

First, To do this part, rst you need to start Kafka and create a topic, named avg,

```
zookeeper-server-start.sh $KAFKA_HOME/config/zookeeper.properties
```

```
kafka-server-start.sh $KAFKA_HOME/config/server.properties
```

```
kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic avg
```

Second, run Cassandra,

```
cassandra -f
```

Third, enter generator folder and run,

```
sbt run
```

(string, array) will be inputted into topic avg as shown below

```
(null,i,0)
(null,t,17)
(null,y,3)
(null,t,7)
(null,v,25)
(null,z,14)
(null,h,9)
(null,x,0)
(null,f,7)
(null,a,25)
```

The last step is to enter sparkstreaming folder and run,

```
sbt run
```

The output is shown as

```
(i,12.410886858194727)
(t,12.506366047745358)
(y,12.58083596214511)
(t,12.506000928320404)
(v,12.495751994292755)
(z,12.440589878516208)
(h,12.450692211797126)
(x,12.574443065264258)
(f,12.46871541115958)
(a,12.531992449391394)
```

Also, if start the cqlsh prompt and print the list of keyspaces in Cassandra as
`$CASSANDRA_HOME/bin/cqlsh`
describe keyspaces;
You can see the table.